

Biostat M280 HW2

Due 05/05

Julia Cheatsheet <https://cheatsheets.quantecon.org/julia-cheatsheet.html>
[\(https://cheatsheets.quantecon.org/julia-cheatsheet.html\)](https://cheatsheets.quantecon.org/julia-cheatsheet.html)

Q1

1) Implement Kinship estimator.

First, we aim to optimize the computation of e_{ij} .

```
In [1]: # variant 1: loops to check correctness (slow? yes)
function e1(X::Matrix{Float64})
    (n,m) = size(X)
    e = zeros(X)
    for i = 1:n, j = 1:n
        for k = 1:m
            e[i,j] += 0.25 * (X[i,k]*X[j,k]+(2-X[i,k])*(2-X[j,k]))
        end
    end
    return e
end

# variant 2: vectorized (fast? no, actually slow)
function e2(X::Matrix{Float64})
    (n,m) = size(X)
    e = zeros(X)
    X2 = (2 - X)
    for i = 1:n, j = 1:n
        xx = (X[i,:]' * X[j,:])[1]
        xx2 = (X2[i,:]' * X2[j,:])[1]
        e[i,j] = 0.25 .* (xx + xx2)
    end
    return e
end

# variant 3: matrix ops (very fast? yes)
function e3(X::Matrix{Float64})
    return 0.25 .* (X * X' + (2 - X) * (2 - X)')
end
```

```
Out[1]: e3 (generic function with 1 method)
```

```
In [2]: srand(1234)
X = rand(0.0:2.0, 10, 10)
e1(X)
string(sum(abs(e1(X) - e2(X)) .> 10e-5)) * " differences"
```

```
Out[2]: "0 differences"
```

```
In [3]: string(sum(abs(e1(X) - e3(X)) .> 10e-5)) * " differences"
```

```
Out[3]: "0 differences"
```

```
In [4]: using BenchmarkTools
srand(280)
X = rand(0.0:2.0, 100, 100)
@benchmark e1(X)
```

```
Out[4]: BenchmarkTools.Trial:
  memory estimate:  78.20 KiB
  allocs estimate:  2
  -----
  minimum time:      3.211 ms (0.00% GC)
  median time:       3.253 ms (0.00% GC)
  mean time:         3.291 ms (0.10% GC)
  maximum time:      5.895 ms (0.00% GC)
  -----
  samples:           1513
  evals/sample:      1
```

```
In [5]: using BenchmarkTools
srand(280)
X = rand(0.0:2.0, 100, 100)
@benchmark e2(X)
```

```
Out[5]: BenchmarkTools.Trial:
  memory estimate:  36.47 MiB
  allocs estimate:  80004
  -----
  minimum time:      12.080 ms (12.84% GC)
  median time:       15.555 ms (22.05% GC)
  mean time:         17.161 ms (21.02% GC)
  maximum time:      93.667 ms (85.58% GC)
  -----
  samples:           291
  evals/sample:      1
```

```
In [6]: using BenchmarkTools
        srand(280)
        X = rand(0.0:2.0, 100, 100)
        @benchmark e3(X)
```

```
Out[6]: BenchmarkTools.Trial:
  memory estimate: 469.22 KiB
  allocs estimate: 12
  -----
  minimum time:      191.685 μs (0.00% GC)
  median time:       259.976 μs (0.00% GC)
  mean time:         301.669 μs (7.23% GC)
  maximum time:      5.129 ms (0.00% GC)
  -----
  samples:           10000
  evals/sample:      1
```

Now we optimize the computation of $\hat{\Phi}$.

```
In [7]: # variant 1: loops (slow? yes)
        function kinship1(X::Matrix{Float64})
            (n,m) = size(X)
            e = e3(X)

            p = zeros(Float64, m)
            for k = 1:m
                p[k] = (1/(2*n))*sum(X[:,k])
            end

            sum_p = 0
            for k = 1:m
                sum_p += (p[k]^2 + (1- p[k])^2)
            end

            Phi = zeros(Float64, (n,n))
            for i = 1:n, j = 1:n
                Phi[i,j] = (e[i,j]-sum_p) / (m - sum_p)
            end

            return Phi
        end
```

```
Out[7]: kinship1 (generic function with 1 method)
```

```
In [8]: # variant 2: avoid loops (faster? yes)
function kinship2(X::Matrix{Float64})
    (n,m) = size(X)
    e = e3(X)
    p = (1/(2*n)) * sum(X, 1)
    sum_p = m + 2*sum(p.^2 - p) # expand p^2 + (1 - p)^2 and simplify
    Phi = (e - sum_p) ./ (m - sum_p)
    return Phi
end
```

```
Out[8]: kinship2 (generic function with 1 method)
```

```
In [9]: using BenchmarkTools
srand(280)
X = rand(0.0:2.0, 1000, 10000)
@benchmark kinship1(X)
```

```
Out[9]: BenchmarkTools.Trial:
  memory estimate:  360.48 MiB
  allocs estimate:  6049505
  -----
  minimum time:      843.138 ms (17.51% GC)
  median time:       958.979 ms (16.49% GC)
  mean time:         1.002 s (23.20% GC)
  maximum time:      1.355 s (38.81% GC)
  -----
  samples:           5
  evals/sample:      1
```

```
In [22]: using BenchmarkTools
srand(280)
X = rand(0.0:2.0, 1000, 10000)
@benchmark kinship2(X)
```

```
Out[22]: BenchmarkTools.Trial:
  memory estimate:  198.67 MiB
  allocs estimate:  27
  -----
  minimum time:      368.366 ms (2.70% GC)
  median time:       506.146 ms (7.29% GC)
  mean time:         500.862 ms (13.87% GC)
  maximum time:      642.811 ms (27.00% GC)
  -----
  samples:           10
  evals/sample:      1
```

2) Benchmark your implementation.

```
In [33]: using BenchmarkTools
          srand(280)
          X = rand(0.0:2.0, 1000, 10000)
          @benchmark kinship2(X)
```

```
Out[33]: BenchmarkTools.Trial:
          memory estimate: 198.67 MiB
          allocs estimate: 27
          -----
          minimum time:      334.318 ms (1.35% GC)
          median time:      413.373 ms (8.16% GC)
          mean time:        464.584 ms (14.96% GC)
          maximum time:     688.434 ms (26.77% GC)
          -----
          samples:           11
          evals/sample:      1
```

We achieve a minimum time of 335ms using 200MiB memory. Note that this benchmark was executed on a juliabox.com instance (possibly shared by multiple virtual environments / users).

```
In [12]: versioninfo()

Julia Version 0.5.1
Commit 6445c82 (2017-03-05 13:25 UTC)
Platform Info:
  OS: Linux (x86_64-pc-linux-gnu)
  CPU: Intel(R) Xeon(R) CPU @ 2.60GHz
  WORD_SIZE: 64
  BLAS: libopenblas (USE64BITINT DYNAMIC_ARCH NO_AFFINITY Sandybridge)
  LAPACK: libopenblas64_
  LIBM: libopenlibm
  LLVM: libLLVM-3.7.1 (ORCJIT, sandybridge)
```

Q2

1) Sherman-Morrison

Show $L^{-1} = R \Leftrightarrow LR = RL = I$.

(1) $LR = I$.

$$\begin{aligned}
LR &= (A + uu^T) \left(A^{-1} - \frac{A^{-1}uu^T A^{-1}}{1 + u^T A^{-1}u} \right) \\
&= AA^{-1} + uv^T A^{-1} - \frac{AA^{-1}uv^T A^{-1} + uv^T A^{-1}uv^T A^{-1}}{1 + v^T A^{-1}u} \\
&= I + uv^T A^{-1} - \frac{uv^T A^{-1} + uv^T A^{-1}uv^T A^{-1}}{1 + v^T A^{-1}u} \\
&= I + uv^T A^{-1} - \frac{u(1 + v^T A^{-1}u)v^T A^{-1}}{1 + v^T A^{-1}u} \\
&= I + uv^T A^{-1} - uv^T A^{-1} \\
&= I.
\end{aligned}$$

(2) $RL = I$.

$$\begin{aligned}
RL &= \left(A^{-1} - \frac{A^{-1}uu^T A^{-1}}{1 + u^T A^{-1}u} \right) (A + uu^T) \\
&= I.
\end{aligned}$$

QED.

2) Woodbury

Apply binomial inverse theorem (3) with $B = I$. QED.

3) Binomial inversion

$$\begin{aligned}
&(A + UBV^T)(A^{-1} - A^{-1}U(B^{-1} + V^T A^{-1}U)^{-1}V^T A^{-1}) \\
&= (A + UBV^T)(A^{-1} - A^{-1}UB(B + V^T A^{-1}UB)^{-1}BV^T A^{-1}) \\
&= (A + UBV^T)(A^{-1}) - ((A + UBV^T)(A^{-1})(UB(B + V^T A^{-1}UB)^{-1}BV^T A^{-1})) \\
&= (A + UBV^T)(A^{-1}) - ((A + UBV^T)(A^{-1}UB)(B + V^T A^{-1}UB)^{-1}BV^T A^{-1}) \\
&= (A + UBV^T)(A^{-1}) - (U(B + BV^T A^{-1}UB))(B + V^T A^{-1}UB)^{-1}BV^T A^{-1} \\
&= I + UBV^T A^{-1} - U(B + BV^T A^{-1}UB)(B + V^T A^{-1}UB)^{-1}BV^T A^{-1} \\
&= I + UBV^T A^{-1} - UBV^T A^{-1} \\
&= I.
\end{aligned}$$

QED.

4) Recall Sylvester's determinant theorem

$$\det(I_m + AB) = \det(I_n + BA).$$

Observe

$$\det(X + AB) = \det(X) \det(I_n + BX^{-1}A).$$

Apply

$$\begin{aligned} \det(A + UV^T) &= \det A \det(I_n + A^{-1}UV^T) \\ &= \det A \det \begin{pmatrix} I & V^T \\ -A^{-1} & I \end{pmatrix} \\ &= \det A \det(I_m + V^T A^{-1}U). \end{aligned}$$

QED.

Q3

We know

$$y_i = x_i^T \beta + z_i^T \gamma + \epsilon_i$$

where $\epsilon \sim N(0, \sigma_0^2)$ and $\gamma \sim N(0_q, \sigma_1^2 I_q)$.

1) Show

$$y \sim N(X\beta, \sigma_0^2 \cdot I + \sigma_1^2 ZZ^T).$$

Note that the sum of normally distributed random variables is normally distributed. Observe

$$\begin{aligned} E(y) &= E(X\beta + Z\gamma + \epsilon) \\ &= X\beta + ZE(\gamma) + E(\epsilon) . \\ &= X\beta \end{aligned}$$

And

$$\begin{aligned} \text{Var}(y) &= \text{Var}(X\beta + Z\gamma + \epsilon) \\ &= \text{Var}(X\beta) + Z\text{Var}(\gamma)Z^T + \text{Var}(\epsilon) + 2\text{Cov}(X\beta, Z\gamma) + 2\text{Cov}(X\beta, \epsilon) + 2\text{Cov}(Z\gamma, \epsilon) \\ &= \sigma_1^2 ZZ^T + \sigma_0^2 I. \end{aligned}$$

QED.

2) Implement function logpdf.

```
In [13]: function logpdf_mvn(y::Vector, Z::Matrix, sigma0, sigma1)
          sigma = sigma0^2 * I + sigma1^2 * Z * Z'
          n = length(y)
          sigmachol = cholfact(sigma)
          - (n/2) * log(2π) - (1/2) * logdet(sigmachol) - (1/2) * sumabs2(sigmachol[L] \ y)
end
```

```
Out[13]: logpdf_mvn (generic function with 1 method)
```

3) Benchmark your result.

```
In [14]: using BenchmarkTools, Distributions

srand(280)

n, q = 2000, 10
Z = randn(n, q)
sigma0, sigma1 = 0.5, 2.0

sigma = sigma1^2 * Z * Z' + sigma0^2 * I
mvn = MvNormal(sigma)
y = rand(mvn)

# check you answer matches that from Distributions.jl
@show logpdf_mvn(y, Z, sigma0, sigma1)
@show logpdf(mvn, y)

# benchmark
@benchmark logpdf_mvn(y, Z, sigma0, sigma1)

logpdf_mvn(y,Z,sigma0,sigma1) = -1571.5736734654183
logpdf(mvn,y) = -1571.5736734654186
```

```
Out[14]: BenchmarkTools.Trial:
  memory estimate: 122.39 MiB
  allocs estimate: 25
  -----
  minimum time:      171.113 ms (23.83% GC)
  median time:       217.344 ms (33.17% GC)
  mean time:         229.807 ms (33.25% GC)
  maximum time:      343.878 ms (38.90% GC)
  -----
  samples:           22
  evals/sample:      1
```



```
In [15]: @benchmark logpdf(mvn, y)
```

```
Out[15]: BenchmarkTools.Trial:
  memory estimate: 15.78 KiB
  allocs estimate: 3
  -----
  minimum time:      9.624 ms (0.00% GC)
  median time:      13.829 ms (0.00% GC)
  mean time:        14.273 ms (0.00% GC)
  maximum time:     40.516 ms (0.00% GC)
  -----
  samples:           349
  evals/sample:      1
```